

# senstoolkit

Sensitivity Analysis Toolkit — User Manual

Version 0.1.0

February 21, 2026

## Contents

<b>1</b>	<b>Introduction</b>	<b>4</b>
<b>2</b>	<b>Installation</b>	<b>4</b>
2.1	Basic Installation . . . . .	4
2.2	Full Installation (Recommended) . . . . .	4
2.3	Optional Dependencies and Feature Gating . . . . .	4
<b>3</b>	<b>Quick Start</b>	<b>4</b>
3.1	Step 1: Create a Parameter File . . . . .	5
3.2	Step 2: Generate a DOE . . . . .	5
3.3	Step 3: Fill in the Response Values . . . . .	5
3.4	Step 4: Run the Analysis . . . . .	5
<b>4</b>	<b>Extending an Existing DOE</b>	<b>5</b>
<b>5</b>	<b>Command-Line Interface Reference</b>	<b>6</b>
5.1	<code>senstoolkit template</code> . . . . .	6
5.2	<code>senstoolkit design</code> . . . . .	6
5.3	<code>senstoolkit extend</code> . . . . .	6
5.4	<code>senstoolkit analyze</code> . . . . .	7
<b>6</b>	<b>Python API Reference</b>	<b>7</b>
6.1	Key Function Signatures . . . . .	7
<b>7</b>	<b>Output Files Reference</b>	<b>7</b>
<b>8</b>	<b>Analysis Methods: Background and Interpretation</b>	<b>8</b>
8.1	Scatter Plot Grid . . . . .	8
8.1.1	What It Does . . . . .	8
8.1.2	Background . . . . .	8
8.1.3	Interpretation . . . . .	8
8.1.4	Output . . . . .	9
8.2	Pearson and Spearman Correlation . . . . .	9
8.2.1	What It Does . . . . .	9
8.2.2	Mathematical Background . . . . .	9
8.2.3	Interpretation . . . . .	9
8.2.4	Output . . . . .	9

8.3	Surrogate Model and Cross-Validation . . . . .	10
8.3.1	What It Does . . . . .	10
8.3.2	Background . . . . .	10
8.3.3	Surrogate Quality Gate . . . . .	10
8.3.4	Output . . . . .	10
8.4	Permutation Importance (Cross-Validated) . . . . .	10
8.4.1	What It Does . . . . .	10
8.4.2	Background . . . . .	10
8.4.3	Interpretation . . . . .	11
8.4.4	Output . . . . .	11
8.5	XGBoost Gain Importance . . . . .	11
8.5.1	What It Does . . . . .	11
8.5.2	Background . . . . .	11
8.5.3	Interpretation . . . . .	11
8.5.4	Output . . . . .	11
8.6	Grouped Permutation Importance . . . . .	12
8.6.1	What It Does . . . . .	12
8.6.2	Background . . . . .	12
8.6.3	Interpretation . . . . .	12
8.6.4	Output . . . . .	12
8.7	Morris Elementary Effects Screening . . . . .	12
8.7.1	What It Does . . . . .	12
8.7.2	Mathematical Background . . . . .	12
8.7.3	Interpretation . . . . .	13
8.7.4	Output . . . . .	13
8.8	Partial Dependence Plots (PDP) and Individual Conditional Expectation (ICE) . . . . .	13
8.8.1	What It Does . . . . .	13
8.8.2	Background . . . . .	13
8.8.3	Interpretation . . . . .	13
8.8.4	Output . . . . .	14
8.9	SHAP (SHapley Additive exPlanations) . . . . .	14
8.9.1	What It Does . . . . .	14
8.9.2	Mathematical Background . . . . .	14
8.9.3	Interpretation . . . . .	14
8.9.4	Output . . . . .	14
8.10	SHAP Interaction Values . . . . .	15
8.10.1	What It Does . . . . .	15
8.10.2	Background . . . . .	15
8.10.3	Interpretation . . . . .	15
8.10.4	Output . . . . .	15
8.11	Sobol Sensitivity Indices . . . . .	15
8.11.1	What It Does . . . . .	15
8.11.2	Mathematical Background . . . . .	15
8.11.3	Interpretation . . . . .	16
8.11.4	Output . . . . .	16
9	Comparing Methods: When to Use What	17
10	Design of Experiments: Background	17
10.1	Sobol Quasi-Random Sequences . . . . .	17
10.2	Linear vs. Log Scale . . . . .	18
10.3	Sample Size Guidelines . . . . .	18

**11 Troubleshooting** **18**

**12 References** **18**

# 1 Introduction

`senstoolkit` is a Python package for performing comprehensive sensitivity analysis on simulation or experimental data. It covers the full workflow from experimental design through multi-method analysis:

1. **Design of Experiments (DOE)** — generate space-filling Sobol quasi-random samples over user-defined parameter ranges.
2. **Multi-method sensitivity analysis** — run 12 complementary analysis methods on the filled DOE, producing CSV tables and publication-ready plots.

The toolkit is fully generic: it works with arbitrary parameter names, arbitrary numbers of parameters, and arbitrary response (output) variables. No domain-specific defaults are hard-coded.

## 2 Installation

### 2.1 Basic Installation

```
cd sensitivity_analysis/  
pip install -e .
```

This installs the core dependencies: `numpy`, `pandas`, `scipy`, `scikit-learn`, and `matplotlib`.

### 2.2 Full Installation (Recommended)

To enable all analysis methods (XGBoost surrogate, SHAP, Sobol via SALib, Morris screening):

```
pip install -e ".[full]"
```

This additionally installs `xgboost`, `shap`, and `SALib`.

### 2.3 Optional Dependencies and Feature Gating

Package	Features enabled	Fallback
<code>xgboost</code>	XGBoost surrogate, gain importance, SHAP	<code>HistGradientBoostingRegressor</code>
<code>shap</code>	SHAP values, SHAP interactions	Skipped
<code>SALib</code>	Sobol indices (SALib), Morris screening	Manual Sobol / skipped

Table 1: Optional dependencies and their fallback behaviour.

## 3 Quick Start

A typical workflow consists of four steps:

1. **Create a parameter template** and edit it to define your parameters.
2. **Generate a DOE** (Design of Experiments).
3. **Run your simulation** to fill in the response columns.
4. **Run the analysis** to obtain sensitivity results.

### 3.1 Step 1: Create a Parameter File

```
senstoolkit template --output my_params.json
```

This creates a JSON file with example parameters. Edit it to define your own:

```
{  
    "temperature": {"min": 300, "max": 800, "scale": "linear"},  
    "pressure": {"min": 0.1, "max": 100, "scale": "log"},  
    "flow_rate": {"min": 0.5, "max": 5.0, "scale": "linear"}  
}
```

Each parameter requires:

- `min`, `max` — the lower and upper bounds.
- `scale` — "linear" (uniform sampling) or "log" (log-uniform sampling; requires `min > 0`).

### 3.2 Step 2: Generate a DOE

```
senstoolkit design --params my_params.json \  
    --n-samples 128 \  
    --seed 42 \  
    --output doe.csv \  
    --response-cols yield efficiency
```

This produces:

- `doe.csv` — a CSV with columns: `id`, parameter columns, and empty response columns (`yield`, `efficiency`).
- `doe.params.json` — a sidecar file storing the parameter bounds, scales, seed, and sample count (needed for extending and analysing).

### 3.3 Step 3: Fill in the Response Values

Open `doe.csv` in a spreadsheet or run your simulation code to populate the `yield` and `efficiency` columns with the corresponding output values for each parameter combination.

### 3.4 Step 4: Run the Analysis

```
senstoolkit analyze --csv doe.csv \  
    --response-cols yield efficiency \  
    --out-dir results/
```

All output files (CSV tables and PNG plots) are written to the `results/` directory.

## 4 Extending an Existing DOE

Sobol sequences are extensible by design. If you need more samples after an initial run, you can extend the DOE without discarding existing simulation results:

```
senstoolkit extend --csv doe.csv --n-new 256
```

This:

- Reads the original seed from the sidecar file (`doe.params.json`).
- Uses `Sobol.fast_forward()` to skip past the existing points.
- Generates 256 new quasi-random samples and appends them to `doe.csv`.
- Preserves all existing rows (including already-filled response values) unchanged.
- Updates the sidecar metadata.

You can also write to a separate file:

```
senstoolkit extend --csv doe.csv --n-new 256 --output doe_extended.csv
```

## 5 Command-Line Interface Reference

### 5.1 `senstoolkit template`

Flag	Default	Description
<code>-output</code>	<code>parameters_template.json</code>	Output file path

### 5.2 `senstoolkit design`

Flag	Default	Description
<code>-params</code>	(required)	Path to parameters JSON file
<code>-n-samples</code>	<code>max(10p, 50)</code>	Number of samples
<code>-output</code>	<code>DOE_&lt;timestamp&gt;.csv</code>	Output CSV path
<code>-seed</code>	None	Random seed for reproducibility
<code>-response-cols</code>	(empty)	Response column names to add

### 5.3 `senstoolkit extend`

Flag	Default	Description
<code>-csv</code>	(required)	Path to the existing DOE CSV
<code>-n-new</code>	(required)	Number of new samples to add
<code>-output</code>	(overwrite existing)	Output CSV path
<code>-seed</code>	(from sidecar)	Override the original seed

## 5.4 senstoolkit analyze

Flag	Default	Description
-csv	(required)	Path to the filled DOE CSV
-response-cols	(required)	Response column name(s)
-out-dir	outputs	Output directory
-seed	0	Random seed
-cv-folds	5	Number of cross-validation folds
-r2-threshold	0.5	Minimum CV $R^2$ for surrogate methods
-no-pdp		Skip PDP/ICE plots
-no-shap		Skip SHAP analysis
-no-group-perm		Skip grouped permutation importance
-no-sobol		Skip Sobol analysis
-no-morris		Skip Morris screening
-no-scatter		Skip scatter plot grid

## 6 Python API Reference

All functions are importable from the top-level package:

```
from senstoolkit import (
    design, extend_design, analyze,
    write_params_template, morris_screening,
    parse_params_json, sobol_sample, apply_scaling, write_doe_csv,
)
```

Submodule imports for advanced use:

```
from senstoolkit.importance import fit_model, cv_permutation_importance
from senstoolkit.sobol import sobol_on_surrogate
from senstoolkit.correlation import correlation_vector, bootstrap_correlation_ci
from senstoolkit.plotting import scatter_grid, shap_interaction_analysis
```

### 6.1 Key Function Signatures

```
def design(params_json, n_samples, out_csv,
           seed=None, prefer_power_of_two=False,
           response_cols=()):

def extend_design(existing_csv, n_new,
                  out_csv=None, seed=None):

def analyze(design_csv, response_cols, out_dir="outputs",
            seed=0, cv_folds=5, perm_repeats=20,
            bootstrap_corr=1000, group_corr_threshold=0.9,
            top_k_pdp=6, do_pdp=True, do_shap=True,
            do_group_perm=True, do_sobol=True,
            sobol_samples=2048, do_morris=True,
            do_scatter=True, r2_threshold=0.5):
```

## 7 Output Files Reference

For each response column <Y>, the analysis produces the following files in the output directory:

File	Description
summary.json	Overall summary: $R^2$ scores, sample counts, surrogate quality
scatter_grid_<Y>.png	Scatter plot grid of each parameter vs. response
correlations_<Y>.csv	Spearman & Pearson correlations with bootstrap 95% CIs
corr_spearman_<Y>.png	Spearman correlation bar chart
corr_pearson_<Y>.png	Pearson correlation bar chart
perm_cv_<Y>.csv	Cross-validated permutation importance (mean $R^2$ drop)
pareto_perm_cv_<Y>.png	Permutation importance bar chart
xgb_gain_<Y>.csv	XGBoost gain importance values
pareto_gain_<Y>.png	Gain importance bar chart
perm_groups_<Y>.csv	Grouped permutation importance
pareto_perm_groups_<Y>.png	Grouped importance bar chart
feature_corr_<Y>.csv	Feature-to-feature correlation matrix
morris_<Y>.csv	Morris $\mu^*$ and $\sigma$ values
morris_mustar_<Y>.png	Morris $\mu^*$ bar chart
morris_sigma_<Y>.png	Morris $\sigma$ bar chart
pdp/pdp_<Y>_<param>.png	PDP/ICE plots for top parameters
shap_bar_<Y>.png	SHAP mean  value  bar chart
shap/shap_dep_<Y>_<param>.png	SHAP dependence plots
shap_interactions_<Y>.csv	SHAP interaction matrix
shap_interactions_<Y>.png	SHAP interaction heatmap
sobol_<Y>.csv	Sobol $S_1$ and $S_T$ indices
sobol_S1_<Y>.png	Sobol first-order bar chart
sobol_ST_<Y>.png	Sobol total-order bar chart
sobol_S2_<Y>.csv	Sobol second-order interaction matrix
sobol_S2_<Y>.png	Sobol $S_2$ interaction heatmap

## 8 Analysis Methods: Background and Interpretation

This section explains each sensitivity analysis method in detail. For each method we describe: (1) what it computes, (2) the mathematical background, (3) the output files, and (4) how to interpret the results.

### 8.1 Scatter Plot Grid

#### 8.1.1 What It Does

For each input parameter  $x_i$ , a scatter plot of  $x_i$  versus the response  $y$  is produced. All subplots are arranged in a single grid image.

#### 8.1.2 Background

This is the simplest and most direct visual method. It makes no assumptions about the relationship between parameters and response. You see the raw data.

#### 8.1.3 Interpretation

- **Clear trend** (upward/downward slope, curve) — the parameter has a visible effect on the response.

- **Funnel shape** (variance changes with  $x_i$ ) — the parameter influences not only the mean but also the variability of the response (heteroscedasticity).
- **Flat cloud** (no pattern) — the parameter has little or no influence.
- **Clusters or gaps** — may indicate discrete regimes, thresholds, or nonlinear effects.

#### 8.1.4 Output

`scatter_grid_<Y>.png`

## 8.2 Pearson and Spearman Correlation

### 8.2.1 What It Does

Computes the Pearson and Spearman rank correlation coefficient between each parameter  $x_i$  and the response  $y$ , along with bootstrap 95% confidence intervals.

### 8.2.2 Mathematical Background

**Pearson correlation ( $r$ )** measures the strength of the *linear* relationship:

$$r_{x_i,y} = \frac{\sum_{k=1}^n (x_{ik} - \bar{x}_i)(y_k - \bar{y})}{\sqrt{\sum_{k=1}^n (x_{ik} - \bar{x}_i)^2} \cdot \sqrt{\sum_{k=1}^n (y_k - \bar{y})^2}} \quad (1)$$

$r = +1$  means perfect positive linear relationship;  $r = -1$  means perfect negative linear relationship;  $r = 0$  means no linear relationship (but there may still be a nonlinear one).

**Spearman correlation ( $\rho$ )** replaces values by their ranks before computing the Pearson formula. It measures the strength of any *monotonic* relationship (linear or not). If  $\rho$  is large but  $r$  is small, the relationship is monotonic but nonlinear (e.g., logarithmic).

**Bootstrap confidence intervals.** The toolkit resamples the data (with replacement) 1000 times by default and computes the correlation on each resample. The 2.5th and 97.5th percentiles of the resampled distribution form the 95% CI. If the CI does not include zero, the correlation is statistically significant at the 5% level.

### 8.2.3 Interpretation

- $|r|$  or  $|\rho| > 0.7$  — strong relationship.
- $0.3 < |r|$  or  $|\rho| \leq 0.7$  — moderate relationship.
- $|r|$  or  $|\rho| \leq 0.3$  — weak or no relationship.
- **Sign** indicates direction: positive means  $y$  increases as  $x_i$  increases.
- If the bootstrap CI includes zero, the correlation is not statistically significant.
- Compare Pearson vs. Spearman: a large difference suggests a nonlinear (but monotonic) effect.

#### 8.2.4 Output

`correlations_<Y>.csv, corr_spearman_<Y>.png, corr_pearson_<Y>.png`

## 8.3 Surrogate Model and Cross-Validation

### 8.3.1 What It Does

Fits a machine learning regression model (XGBoost or HistGradientBoosting) to approximate the mapping  $\mathbf{x} \mapsto y$ . The model quality is assessed via  $k$ -fold cross-validation  $R^2$ .

### 8.3.2 Background

Many sensitivity analysis methods (Sobol, PDP, SHAP, Morris as implemented here) require evaluating the response at arbitrary parameter combinations. Since only the DOE samples have actual simulation results, a *surrogate model* (also called emulator or metamodel) is trained to predict  $y$  from  $\mathbf{x}$ .

The surrogate is an XGBoost gradient-boosted tree ensemble (400 trees, learning rate 0.05, max depth 6). If XGBoost is not installed, scikit-learn's `HistGradientBoostingRegressor` is used as a fallback.

**Cross-validation  $R^2$ .** The data is split into  $k$  folds (default  $k = 5$ ). For each fold, the model is trained on  $k - 1$  folds and evaluated on the held-out fold. The  $R^2$  score measures how much variance is explained:

$$R^2 = 1 - \frac{\sum_i (y_i - \hat{y}_i)^2}{\sum_i (y_i - \bar{y})^2} \quad (2)$$

$R^2 = 1$  means perfect prediction;  $R^2 = 0$  means the model is no better than predicting the mean;  $R^2 < 0$  means the model is worse than the mean.

### 8.3.3 Surrogate Quality Gate

If the mean CV  $R^2$  falls below the threshold (default 0.5), the toolkit prints a warning and **skips all surrogate-dependent methods** (PDP, SHAP, Sobol, Morris) for that response. This prevents misleading results from a poor surrogate.

#### What to do if the quality gate triggers:

- Add more samples (`senstoolkit extend`).
- Check for outliers or data errors in the response column.
- The non-surrogate methods (correlation, scatter, permutation importance) still run and may provide useful insights.

### 8.3.4 Output

`summary.json` (contains `cv_r2_mean`, `cv_r2_std`, `train_r2`, `surrogate_ok`)

## 8.4 Permutation Importance (Cross-Validated)

### 8.4.1 What It Does

Measures how much the model's predictive performance ( $R^2$ ) drops when each parameter's values are randomly shuffled.

### 8.4.2 Background

Permutation importance (Breiman, 2001) works as follows:

1. Train the model on the training fold and compute the baseline  $R^2$  on the validation fold.

2. For each parameter  $x_i$ : randomly shuffle (permute) the values of  $x_i$  in the validation set, re-predict, and measure the new  $R^2$ .
3. The *importance* of  $x_i$  is the drop in  $R^2$ :

$$\text{Importance}(x_i) = R_{\text{baseline}}^2 - R_{\text{permuted}}^2 \quad (3)$$

4. Repeat across all CV folds and multiple permutation repeats; report the mean and standard deviation.

This is **model-agnostic**: it measures the importance of a feature to the model's predictions, regardless of the model type.

#### 8.4.3 Interpretation

- **Large positive value**: the parameter is important — shuffling it destroys predictive power.
- **Near zero**: the parameter contributes little; the model can predict almost as well without it.
- **Negative value** (rare): shuffling actually improved the score, which typically indicates noise or overfitting.
- The standard deviation across folds indicates how stable the importance estimate is.

#### 8.4.4 Output

`perm_cv_<Y>.csv`, `pareto_perm_cv_<Y>.png`

### 8.5 XGBoost Gain Importance

#### 8.5.1 What It Does

Reports the total reduction in the loss function (“gain”) attributable to each parameter across all splits in all trees of the XGBoost ensemble.

#### 8.5.2 Background

Each time a tree node splits on parameter  $x_i$ , the split produces a gain: the improvement in the loss function. The *gain importance* of  $x_i$  is the sum (or average) of these gains over all trees and all splits involving  $x_i$ .

This is an **internal metric of the XGBoost model**. It is fast to compute (no re-evaluation needed) but can be biased towards high-cardinality or correlated features.

#### 8.5.3 Interpretation

- Higher gain = more important for the model's predictions.
- Compare with permutation importance: if gain is high but permutation importance is low, the feature may be redundant (another correlated feature carries the same information).
- Only available when XGBoost is the surrogate model.

#### 8.5.4 Output

`xgb_gain_<Y>.csv`, `pareto_gain_<Y>.png`

## 8.6 Grouped Permutation Importance

### 8.6.1 What It Does

Identifies groups of highly correlated parameters (using hierarchical clustering at a threshold of  $|\text{corr}| \geq 0.9$ ) and permutes entire groups together to assess their joint importance.

### 8.6.2 Background

When two parameters  $x_i$  and  $x_j$  are highly correlated, permuting  $x_i$  alone may not cause a large  $R^2$  drop because  $x_j$  carries similar information. This leads to *underestimation* of their individual importances.

Grouped permutation importance solves this by:

1. Computing the absolute correlation matrix  $|C|$  among all parameters.
2. Using agglomerative hierarchical clustering (complete linkage) with distance  $d = 1 - |C|$  to form groups of correlated parameters.
3. Permuting all parameters in a group *simultaneously* and measuring the joint  $R^2$  drop.

### 8.6.3 Interpretation

- A group with a large importance value means those correlated parameters are jointly important.
- If a group has high grouped importance but each member has low individual importance, the effect is shared among the correlated parameters.
- The `feature_corr_<Y>.csv` file lets you inspect which parameters are correlated.

### 8.6.4 Output

`perm_groups_<Y>.csv`, `pareto_perm_groups_<Y>.png`, `feature_corr_<Y>.csv`

## 8.7 Morris Elementary Effects Screening

### 8.7.1 What It Does

Computes the Morris Elementary Effects (EE) to screen parameters into three categories: negligible, linear/additive, and nonlinear/interactive.

### 8.7.2 Mathematical Background

The Morris method (Morris, 1991) works in a discretized parameter space with  $p$  levels. For each parameter  $x_i$ , an *elementary effect* is:

$$\text{EE}_i = \frac{f(x_1, \dots, x_i + \Delta, \dots, x_d) - f(x_1, \dots, x_i, \dots, x_d)}{\Delta} \quad (4)$$

where  $\Delta$  is a fixed step size. Multiple elementary effects are computed along random *trajectories* through the parameter space.

Two summary statistics are reported:

- $\mu_i^* = \frac{1}{r} \sum_{k=1}^r |\text{EE}_i^{(k)}|$  — the mean of the *absolute* elementary effects. Measures overall influence.

- $\sigma_i = \text{std}(\text{EE}_i^{(1)}, \dots, \text{EE}_i^{(r)})$  — the standard deviation of the elementary effects. Measures nonlinearity and/or interactions.

### 8.7.3 Interpretation

- **High  $\mu^*$ , low  $\sigma$ :** the parameter has a strong, mostly linear/additive effect.
- **High  $\mu^*$ , high  $\sigma$ :** the parameter has a strong effect that is nonlinear and/or interacts with other parameters.
- **Low  $\mu^*$ , low  $\sigma$ :** the parameter has negligible influence and can potentially be fixed.
- Morris is a *screening method*: it efficiently identifies which parameters matter most, especially useful when the number of parameters is large ( $d > 10$ ) and full Sobol analysis would be too expensive.

### 8.7.4 Output

`morris_<Y>.csv`, `morris_mustar_<Y>.png`, `morris_sigma_<Y>.png`

**Requires:** SALib package and a sidecar parameter file.

## 8.8 Partial Dependence Plots (PDP) and Individual Conditional Expectation (ICE)

### 8.8.1 What It Does

Shows how the predicted response changes as one parameter varies, while averaging over (PDP) or showing individual traces for (ICE) the other parameters.

### 8.8.2 Background

The *partial dependence function* for parameter  $x_i$  is:

$$\hat{f}_i(x_i) = \frac{1}{n} \sum_{k=1}^n \hat{f}(x_i, \mathbf{x}_{-i}^{(k)}) \quad (5)$$

where  $\mathbf{x}_{-i}^{(k)}$  denotes all other parameters from the  $k$ -th sample. In other words: for each value of  $x_i$  on a grid, we predict  $y$  for every sample (keeping their other parameter values), then average.

**ICE curves** show the individual predictions  $\hat{f}(x_i, \mathbf{x}_{-i}^{(k)})$  without averaging, so you can see whether the effect of  $x_i$  is the same for all samples or varies (indicating interactions).

### 8.8.3 Interpretation

- **Steep PDP curve:** the parameter has a strong effect on the response.
- **Flat PDP curve:** the parameter has little effect.
- **Nonlinear PDP curve:** the effect is nonlinear (e.g., saturation, threshold).
- **ICE curves that diverge** (spread apart): the effect of  $x_i$  depends on the values of other parameters — an indication of *interactions*.
- **ICE curves that are parallel:** no interactions; the effect of  $x_i$  is the same regardless of other parameter values.

PDPs are generated for the top- $k$  most important parameters (by permutation importance).

#### 8.8.4 Output

pdp/pdp\_<Y>\_<param>.png (one plot per top parameter)

### 8.9 SHAP (SHapley Additive exPlanations)

#### 8.9.1 What It Does

Computes SHAP values for each sample and each parameter, quantifying each parameter's contribution to the prediction.

#### 8.9.2 Mathematical Background

SHAP values are based on Shapley values from cooperative game theory (Shapley, 1953; Lundberg & Lee, 2017). For each sample  $k$  and parameter  $i$ , the SHAP value  $\phi_i^{(k)}$  satisfies:

$$\hat{f}(\mathbf{x}^{(k)}) = \phi_0 + \sum_{i=1}^d \phi_i^{(k)} \quad (6)$$

where  $\phi_0$  is the average prediction. The SHAP value  $\phi_i^{(k)}$  is the parameter's *marginal contribution*, averaged over all possible orderings of parameters.

senstoolkit uses TreeExplainer for tree-based models, which computes exact SHAP values in polynomial time.

**Global importance.** The mean absolute SHAP value  $|\bar{\phi}_i|$  across all samples measures the parameter's global importance.

**Dependence plots.** Plotting  $\phi_i^{(k)}$  vs.  $x_i^{(k)}$  reveals how the parameter's effect varies across its range.

#### 8.9.3 Interpretation

- **Mean  $|\phi_i|$  bar chart:** ranks parameters by global importance (analogous to permutation importance but exact for the model).
- **Dependence plot with a clear trend:** the parameter has a consistent directional effect.
- **Dependence plot with vertical spread at each  $x_i$  value:** interactions with other parameters.
- **Red dashed lines** at  $\pm 0.5\sigma_y$ : SHAP values exceeding half the response standard deviation indicate a practically significant contribution.
- SHAP values sum to the prediction, so they are *additive* and directly comparable across parameters.

#### 8.9.4 Output

shap\_bar\_<Y>.png, shap/shap\_dep\_<Y>\_<param>.png

**Requires:** shap and xgboost packages.

## 8.10 SHAP Interaction Values

### 8.10.1 What It Does

Computes the pairwise SHAP interaction values, producing a  $d \times d$  matrix showing how pairs of parameters jointly influence the prediction.

### 8.10.2 Background

SHAP interaction values decompose the SHAP value into main effects and pairwise interactions:

$$\phi_i^{(k)} = \phi_{ii}^{(k)} + \sum_{j \neq i} \phi_{ij}^{(k)} \quad (7)$$

The diagonal  $\phi_{ii}$  captures the main effect of parameter  $i$ ; the off-diagonal  $\phi_{ij}$  captures the interaction between parameters  $i$  and  $j$ .

### 8.10.3 Interpretation

- **Large diagonal values:** the parameter has a strong main effect.
- **Large off-diagonal values:** the two parameters interact — the effect of one depends on the value of the other.
- The heatmap makes it easy to spot which pairs interact most strongly.
- Useful for identifying which parameter combinations should be explored further (e.g., in 2D PDP or targeted experiments).

### 8.10.4 Output

`shap_interactions_<Y>.csv`, `shap_interactions_<Y>.png`

**Requires:** `shap` and `xgboost` packages.

## 8.11 Sobol Sensitivity Indices

### 8.11.1 What It Does

Decomposes the total variance of the response into contributions from individual parameters (first-order,  $S_1$ ), total effects including all interactions (total-order,  $S_T$ ), and pairwise interactions (second-order,  $S_2$ ).

### 8.11.2 Mathematical Background

Sobol indices (Sobol', 1993) are based on the ANOVA-like decomposition of a function:

$$f(\mathbf{x}) = f_0 + \sum_i f_i(x_i) + \sum_{i < j} f_{ij}(x_i, x_j) + \dots \quad (8)$$

**First-order index  $S_{1,i}$ :**

$$S_{1,i} = \frac{\text{Var}_{x_i} [\mathbb{E}_{\mathbf{x} \sim i} (y | x_i)]}{\text{Var}(y)} \quad (9)$$

This is the fraction of total variance explained by  $x_i$  alone, not considering interactions.

**Total-order index**  $S_{T,i}$ :

$$S_{T,i} = 1 - \frac{\text{Var}_{\mathbf{x}_{\sim i}}[\mathbb{E}_{x_i}(y | \mathbf{x}_{\sim i})]}{\text{Var}(y)} \quad (10)$$

This is the fraction of total variance explained by  $x_i$  and all its interactions with other parameters.

**Second-order index**  $S_{2,ij}$ :

$$S_{2,ij} = \frac{\text{Var}_{x_i, x_j}[\mathbb{E}_{\mathbf{x}_{\sim ij}}(y | x_i, x_j)]}{\text{Var}(y)} - S_{1,i} - S_{1,j} \quad (11)$$

This captures the variance explained by the *interaction* between  $x_i$  and  $x_j$  specifically.

**Estimation.** The toolkit uses the Saltelli (2002, 2010) estimator for  $S_1$  and the Jansen (1999) estimator for  $S_T$ . If SALib is installed, it is used for both  $S_1/S_T$  and  $S_2$ ; otherwise a manual implementation provides  $S_1$  and  $S_T$  (with  $S_2 = 0$  as a placeholder).

### 8.11.3 Interpretation

- $S_{1,i} \approx S_{T,i}$ : the parameter acts mainly alone (no significant interactions).
- $S_{T,i} \gg S_{1,i}$ : a large part of the parameter's influence comes through interactions with other parameters.
- $\sum_i S_{1,i} \approx 1$ : the model is approximately additive (no interactions).
- $\sum_i S_{1,i} \ll 1$ : significant interactions exist.
- $S_{1,i} < 0.05$ : the parameter has negligible direct influence.
- $S_{T,i} < 0.05$ : the parameter has negligible total influence (including interactions) and can be fixed.
- $S_{2,ij}$  identifies specific interacting pairs.

**Diagnostics.** If  $\sum S_{1,i} > 1.05$  or any  $S_T < S_1$ , the toolkit prints a warning. This usually means the sample count is too low; increase `sobol_samples` or add more DOE points.

### 8.11.4 Output

`sobol_<Y>.csv`, `sobol_S1_<Y>.png`, `sobol_ST_<Y>.png`, `sobol_S2_<Y>.csv`, `sobol_S2_<Y>.png`

## 9 Comparing Methods: When to Use What

Method	Needs surrogate?	Captures nonlinearity?	Captures interactions?	Best for
Scatter grid	No	Visual	Visual	Initial exploration
Pearson corr.	No	No	No	Linear screening
Spearman corr.	No	Monotonic	No	Monotonic screening
Perm. importance	Yes (CV)	Yes	Indirectly	Overall ranking
XGBoost gain	Yes	Yes	No	Quick model-internal check
Grouped perm.	Yes (CV)	Yes	Yes (groups)	Correlated parameters
Morris screening	Yes	Yes	Yes ( $\sigma$ )	Efficient screening ( $d > 10$ )
PDP/ICE	Yes	Yes	Yes (ICE)	Shape of effect
SHAP	Yes	Yes	Yes (dependence)	Per-sample attribution
SHAP interactions	Yes	Yes	Yes (pairwise)	Interaction identification
Sobol $S_1/S_T$	Yes	Yes	Yes ( $S_T - S_1$ )	Variance decomposition
Sobol $S_2$	Yes	Yes	Yes (pairwise)	Interaction quantification

Table 3: Comparison of sensitivity analysis methods.

### Recommended workflow:

1. Start with scatter plots and correlations for a quick visual overview.
2. Check the surrogate  $R^2$ . If it is high ( $> 0.8$ ), all surrogate-based methods are reliable.
3. Use permutation importance to rank parameters.
4. Use PDP/ICE to understand the *shape* of each important parameter's effect.
5. Use Sobol indices for a rigorous variance decomposition.
6. Use SHAP interactions or Sobol  $S_2$  to identify interacting parameter pairs.
7. Use Morris screening as a first pass when you have many parameters ( $> 10$ ).

## 10 Design of Experiments: Background

### 10.1 Sobol Quasi-Random Sequences

A Sobol sequence is a *low-discrepancy* (quasi-random) sequence that fills the parameter space more uniformly than pseudo-random sampling. Key properties:

- **Space-filling:** Sobol sequences avoid the clumping and gaps typical of pseudo-random sampling, especially in high dimensions.
- **Deterministic:** given the same seed and dimensionality, the sequence is exactly reproducible.
- **Extensible:** new points can be added without moving existing ones (used by `extend`).

- **Low discrepancy:** the deviation from uniform coverage decreases as  $O((\log n)^d/n)$ , much faster than Monte Carlo's  $O(1/\sqrt{n})$ .

## 10.2 Linear vs. Log Scale

- **Linear:** samples are drawn uniformly between `min` and `max`. Use for parameters where equal absolute changes are equally important (e.g., temperature from 300 to 800 K).
- **Log:** samples are drawn uniformly in log-space, i.e.,  $x = \exp(\text{Uniform}(\ln(\min), \ln(\max)))$ . Use for parameters that span multiple orders of magnitude (e.g., diffusion coefficient from  $10^{-3}$  to  $10^1$ ).

## 10.3 Sample Size Guidelines

- Minimum:  $10 \times d$  samples (where  $d$  is the number of parameters).
- Recommended:  $50\text{--}200 \times d$  for reliable Sobol index estimation.
- The toolkit defaults to  $\max(10d, 50)$  if no sample count is specified.
- Use `senstoolkit extend` to add more samples iteratively.

## 11 Troubleshooting

`surrogate_ok: false`

The surrogate model's CV  $R^2$  is below the threshold. Surrogate-based methods (Sobol, PDP, SHAP, Morris) are skipped. Solutions: add more samples, check for data errors, or lower the threshold via `-r2-threshold`.

**Sobol**  $\sum S_1 > 1$  Insufficient samples for reliable Sobol estimation. Increase `-n-samples` or `sobel_samples`.

**SHAP/Morris skipped**

The required optional packages (`shap`, `SALib`) are not installed. Install with `pip install senstoolkit[full]`.

**Extend fails: “No seed found”**

The sidecar file was generated before the seed-tracking feature was added. Provide `-seed` manually with the same seed used for the original design.

**No sidecar file found**

Sobol and Morris analyses require parameter bounds from the sidecar `.params.json` file. Re-generate the DOE with `senstoolkit design`, or create the sidecar file manually.

## 12 References

- Breiman, L. (2001). Random Forests. *Machine Learning*, 45, 5–32.
- Jansen, M.J.W. (1999). Analysis of variance designs for model output. *Computer Physics Communications*, 117, 35–43.
- Lundberg, S.M. & Lee, S.-I. (2017). A unified approach to interpreting model predictions. *NeurIPS*.
- Morris, M.D. (1991). Factorial sampling plans for preliminary computational experiments. *Technometrics*, 33(2), 161–174.

- Saltelli, A. (2002). Making best use of model evaluations to compute sensitivity indices. *Computer Physics Communications*, 145, 280–297.
- Saltelli, A. et al. (2010). Variance based sensitivity analysis of model output. *Computer Physics Communications*, 181, 259–270.
- Shapley, L.S. (1953). A value for  $n$ -person games. *Contributions to the Theory of Games*, 2, 307–317.
- Sobol', I.M. (1993). Sensitivity estimates for nonlinear mathematical models. *Mathematical Modelling and Computational Experiments*, 1, 407–414.